

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: from sklearn import preprocessing
from sklearn.model_selection import train_test_split
```

```
In [3]: from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import r2_score
```

```
In [4]: data = pd.read_csv(r"C:\Users\DELL\Desktop\FSDS\ML\22nd- 11, 12, scaling\lasso,
data.head()
```

```
Out[4]:
```

	mpg	cyl	disp	hp	wt	acc	yr	origin	car_type	car_name
0	18.0	8	307.0	130	3504	12.0	70	1	0	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70	1	0	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70	1	0	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70	1	0	amc rebel sst
4	17.0	8	302.0	140	3449	10.5	70	1	0	ford torino

```
In [5]: data = data.drop(['car_name'], axis=1)
```

```
In [6]: data['origin'] = data['origin'].replace({1: 'america', 2: 'europe', 3: 'asia'})
data= pd.get_dummies(data, columns = ['origin'], dtype=int)
data= data.replace('?', np.nan)
```

```
In [7]: data=data.apply(pd.to_numeric, errors='ignore')
numeric_cols=data.select_dtypes(include=[np.number]).columns
data[numeric_cols]= data[numeric_cols].apply(lambda x: x.fillna(x.median()))
```

C:\Users\DELL\AppData\Local\Temp\ipykernel\_5376\202028991.py:1: FutureWarning: errors='ignore' is deprecated and will raise in a future version. Use to\_numeric without passing `errors` and catch exceptions explicitly instead

```
data=data.apply(pd.to_numeric, errors='ignore')
```

```
In [8]: data.head()
```

```
Out[8]:
```

	mpg	cyl	disp	hp	wt	acc	yr	car_type	origin_america	origin_asia	origin_europe
0	18.0	8	307.0	130.0	3504	12.0	70	0	1	0	
1	15.0	8	350.0	165.0	3693	11.5	70	0	1	0	
2	18.0	8	318.0	150.0	3436	11.0	70	0	1	0	
3	16.0	8	304.0	150.0	3433	12.0	70	0	1	0	
4	17.0	8	302.0	140.0	3449	10.5	70	0	1	0	

## Model Buliding

```
In [9]: X=data.drop(['mpg'],axis=1)
        y=data[['mpg']]
```

```
In [10]: X_s=preprocessing.scale(X)
        X_columns=X.columns
        X_s=pd.DataFrame(X_s, columns= X_columns)
        y_s=preprocessing.scale(y)
        y_columns=y.columns
        y_s=pd.DataFrame(y_s, columns= y_columns)
```

```
In [13]: X_train, X_test, y_train,y_test = train_test_split(X_s, y_s, test_size = 0.30, r
        X_train.shape
```

```
Out[13]: (278, 10)
```

## Simple Linear Regression

```
In [14]: regression_model = LinearRegression()
        regression_model.fit(X_train, y_train)

        for idx, col_name in enumerate(X_train.columns):
            print('The coefficient for {} is {}'.format(col_name, regression_model.coef_

        intercept = regression_model.intercept_[0]
        print('The intercept is {}'.format(intercept))
```

```
The coefficient for cyl is 0.321022385691611
The coefficient for disp is 0.32483430918483897
The coefficient for hp is -0.22916950059437569
The coefficient for wt is -0.7112101905072298
The coefficient for acc is 0.014713682764191237
The coefficient for yr is 0.3755811949510748
The coefficient for car_type is 0.3814769484233099
The coefficient for origin_america is -0.07472247547584178
The coefficient for origin_asia is 0.044515252035677896
The coefficient for origin_europe is 0.04834854953945386
The intercept is 0.019284116103639764
```

## Regularized Ridge Regression

```
In [15]: ridge_model = Ridge(alpha = 0.3)
        ridge_model.fit(X_train, y_train)

        print('Ridge model coef: {}'.format(ridge_model.coef_))
```

```
Ridge model coef: [[ 0.31649043  0.31320707 -0.22876025 -0.70109447  0.01295851
 0.37447352
 0.37725608 -0.07423624  0.04441039  0.04784031]]
```

## Regularized Lasso Regression

```
In [16]: lasso_model = Lasso(alpha = 0.1)
        lasso_model.fit(X_train, y_train)

        print('Lasso model coef: {}'.format(lasso_model.coef_))
```

```
Lasso model coef: [-0.          -0.          -0.01690287 -0.51890013  0.
0.28138241
 0.1278489  -0.01642647  0.          0.          ]
```

### Score Comparision

```
In [17]: #Simple Linear Model
print(regression_model.score(X_train, y_train))
print(regression_model.score(X_test, y_test))

print('*****')
#Ridge
print(ridge_model.score(X_train, y_train))
print(ridge_model.score(X_test, y_test))

print('*****')
#Lasso
print(lasso_model.score(X_train, y_train))
print(lasso_model.score(X_test, y_test))
```

```
0.8343770256960538
0.8513421387780066
*****
0.8343617931312616
0.8518882171608506
*****
0.7938010766228453
0.8375229615977083
```

### Model Parameter Tuning

```
In [18]: data_train_test = pd.concat([X_train, y_train], axis =1)
data_train_test.head()
```

```
Out[18]:
```

	cyl	displ	hp	wt	acc	yr	car_type	origin_i
<b>350</b>	-0.856321	-0.849116	-1.081977	-0.893172	-0.242570	1.351199	0.941412	0
<b>59</b>	-0.856321	-0.925936	-1.317736	-0.847061	2.879909	-1.085858	0.941412	-1
<b>120</b>	-0.856321	-0.695475	0.201600	-0.121101	-0.024722	-0.815074	0.941412	-1
<b>12</b>	1.498191	1.983643	1.197027	0.934732	-2.203196	-1.627426	-1.062235	0
<b>349</b>	-0.856321	-0.983552	-0.951000	-1.165111	0.156817	1.351199	0.941412	-1

```
In [19]: import statsmodels.formula.api as smf
ols1 = smf.ols(formula = 'mpg ~ cyl+displ+hp+wt+acc+yr+car_type+origin_america+or
ols1.params
```

```
Out[19]: Intercept      0.019284  
        cyl           0.321022  
        disp          0.324834  
        hp            -0.229170  
        wt            -0.711210  
        acc           0.014714  
        yr            0.375581  
        car_type       0.381477  
        origin_america -0.074722  
        origin_europe  0.048349  
        origin_asia    0.044515  
        dtype: float64
```

```
In [20]: print(ols1.summary())
```

## OLS Regression Results

=====						
Dep. Variable:	mpg	R-squared:	0.834			
Model:	OLS	Adj. R-squared:	0.829			
Method:	Least Squares	F-statistic:	150.0			
Date:	Thu, 21 Aug 2025	Prob (F-statistic):	3.12e-99			
Time:	15:53:29	Log-Likelihood:	-146.89			
No. Observations:	278	AIC:	313.8			
Df Residuals:	268	BIC:	350.1			
Df Model:	9					
Covariance Type:	nonrobust					
=====						
=						
	coef	std err	t	P> t	[0.025	0.97
-----						
5]						
-----						
-						
Intercept	0.0193	0.025	0.765	0.445	-0.030	0.06
9						
cyl	0.3210	0.112	2.856	0.005	0.100	0.54
2						
disp	0.3248	0.128	2.544	0.012	0.073	0.57
6						
hp	-0.2292	0.079	-2.915	0.004	-0.384	-0.07
4						
wt	-0.7112	0.088	-8.118	0.000	-0.884	-0.53
9						
acc	0.0147	0.039	0.373	0.709	-0.063	0.09
2						
yr	0.3756	0.029	13.088	0.000	0.319	0.43
2						
car_type	0.3815	0.067	5.728	0.000	0.250	0.51
3						
origin_america	-0.0747	0.020	-3.723	0.000	-0.114	-0.03
5						
origin_europe	0.0483	0.021	2.270	0.024	0.006	0.09
0						
origin_asia	0.0445	0.020	2.175	0.031	0.004	0.08
5						
=====						
Omnibus:	22.678	Durbin-Watson:	2.105			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	36.139			
Skew:	0.513	Prob(JB):	1.42e-08			
Kurtosis:	4.438	Cond. No.	1.59e+16			
=====						

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The smallest eigenvalue is 6.14e-30. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.


```
In [21]: mse = np.mean((regression_model.predict(X_test)-y_test)**2)
import math
rmse = math.sqrt(mse)
print('Root Mean Squared Error: {}'.format(rmse))
```

Root Mean Squared Error: 0.37766934254087847


```
In [25]: fig = plt.figure(figsize=(10,8))
sns.residplot(x= X_test['hp'], y= y_test['mpg'], color='green', lowess=True )
plt.show()

fig = plt.figure(figsize=(10,8))
sns.residplot(x= X_test['acc'], y= y_test['mpg'], color='green', lowess=True )
plt.show()
```





```
In [26]: y_pred = regression_model.predict(X_test)
plt.scatter(y_test['mpg'], y_pred)
plt.show()
```



```
In [ ]:
```

In [ ]: