

Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

Read Dataset

```
In [2]: df = pd.read_csv(r'C:\Users\DELL\Downloads\adult.csv', encoding='latin-1')
```

```
In [3]: df
```

```
Out[3]:
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation
--	-----	-----------	--------	-----------	---------------	----------------	------------

0	90	?	77053	HS-grad	9	Widowed	?
---	----	---	-------	---------	---	---------	---

1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial
---	----	---------	--------	---------	---	---------	-----------------

2	66	?	186061	Some-college	10	Widowed	?
---	----	---	--------	--------------	----	---------	---

3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct
---	----	---------	--------	---------	---	----------	-------------------

4	41	Private	264663	Some-college	10	Separated	Prof-specialty
---	----	---------	--------	--------------	----	-----------	----------------

...
-----	-----	-----	-----	-----	-----	-----	-----

32556	22	Private	310152	Some-college	10	Never-married	Protective-serv
-------	----	---------	--------	--------------	----	---------------	-----------------

32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support
-------	----	---------	--------	------------	----	--------------------	--------------

32558	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct
-------	----	---------	--------	---------	---	--------------------	-------------------

32559	58	Private	151910	HS-grad	9	Widowed	Adm-clerical
-------	----	---------	--------	---------	---	---------	--------------

32560	22	Private	201490	HS-grad	9	Never-married	Adm-clerical
-------	----	---------	--------	---------	---	---------------	--------------

32561 rows × 15 columns



Exploratory Data Analysis

Check shape of dataset

In [4]: `df.shape`


Out[4]: (32561, 15)

Preview dataset

In [5]: `df.head()`

Out[5]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relati
0	90	?	77053	HS-grad	9	Widowed	?	
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	
2	66	?	186061	Some-college	10	Widowed	?	Unr
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unr
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Ow



View summary of dataframe

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education.num         32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital.gain          32561 non-null  int64
11  capital.loss          32561 non-null  int64
12  hours.per.week        32561 non-null  int64
13  native.country        32561 non-null  object
14  income                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

In [8]: `df[df == '?'] = np.nan`

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             30725 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education.num         32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation            30718 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital.gain           32561 non-null  int64
11  capital.loss           32561 non-null  int64
12  hours.per.week         32561 non-null  int64
13  native.country        31978 non-null  object
14  income                 32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

Impute missing values with mode

```
In [10]: for col in ['workclass', 'occupation', 'native.country']:
         df[col].fillna(df[col].mode()[0], inplace=True)
```

Check again for missing values

```
In [12]: df.isnull().sum()
```

```
Out[12]: age                0
workclass                0
fnlwgt                  0
education               0
education.num           0
marital.status          0
occupation              0
relationship            0
race                   0
sex                    0
capital.gain            0
capital.loss            0
hours.per.week          0
native.country          0
income                  0
dtype: int64
```

Setting feature vector and target variable


```
In [14]: X = df.drop(['income'], axis=1)

         y = df['income']
```

```
In [15]: X.head()
```

Out[15]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relati
0	90	Private	77053	HS-grad	9	Widowed	Prof-specialty	
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	
2	66	Private	186061	Some-college	10	Widowed	Prof-specialty	Unr
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unr
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Ow



Split data into separate training and test set

```
In [16]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=42)
```

Feature Engineering

Encode categorical variables

```
In [17]: from sklearn import preprocessing

categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship']
for feature in categorical:
    le = preprocessing.LabelEncoder()
    X_train[feature] = le.fit_transform(X_train[feature])
    X_test[feature] = le.transform(X_test[feature])
```

Feature Scaling

```
In [18]: from sklearn.preprocessing import StandardScaler


scaler = StandardScaler()

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
```

```
In [19]: X_train.head()
```

Out[19]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation
0	0.101484	2.600478	-1.494279	-0.332263	1.133894	-0.402341	-0.782234
1	0.028248	-1.884720	0.438778	0.184396	-0.423425	-0.402341	-0.026690
2	0.247956	-0.090641	0.045292	1.217715	-0.034095	0.926666	-0.782234
3	-0.850587	-1.884720	0.793152	0.184396	-0.423425	0.926666	-0.530380
4	-0.044989	-2.781760	-0.853275	0.442726	1.523223	-0.402341	-0.782234



Logistic Regression model with all features

```
In [20]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred = logreg.predict(X_test)

print('Logistic Regression accuracy score with all the features: {0:0.4f}'.format(accuracy_score(y_test, y_pred)))
```

Logistic Regression accuracy score with all the features: 0.8218

Logistic Regression with PCA

```
In [21]: from sklearn.decomposition import PCA
pca = PCA()
X_train = pca.fit_transform(X_train)
pca.explained_variance_ratio_
```

```
Out[21]: array([0.14757168, 0.10182915, 0.08147199, 0.07880174, 0.07463545,
0.07274281, 0.07009602, 0.06750902, 0.0647268 , 0.06131155,
0.06084207, 0.04839584, 0.04265038, 0.02741548])
```

Logistic Regression with first 13 features

```
In [22]: X = df.drop(['income', 'native.country'], axis=1)
y = df['income']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=42)

categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship']
for feature in categorical:
    le = preprocessing.LabelEncoder()
    X_train[feature] = le.fit_transform(X_train[feature])
    X_test[feature] = le.transform(X_test[feature])

scaler = preprocessing.StandardScaler()
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)

logreg = LogisticRegression()
logreg.fit(X_train, y_train)
```

```
y_pred = logreg.predict(X_test)

print('Logistic Regression accuracy score with the first 13 features: {0:0.4f}').
```

Logistic Regression accuracy score with the first 13 features: 0.8213

Logistic Regression with first 12 features

```
In [23]: X = df.drop(['income', 'native.country', 'hours.per.week'], axis=1)
y = df['income']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=42)

categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship']
for feature in categorical:
    le = preprocessing.LabelEncoder()
    X_train[feature] = le.fit_transform(X_train[feature])
    X_test[feature] = le.transform(X_test[feature])

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)

logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred = logreg.predict(X_test)

print('Logistic Regression accuracy score with the first 12 features: {0:0.4f}').
```

Logistic Regression accuracy score with the first 12 features: 0.8227

Logistic Regression with first 11 features

```
In [24]: X = df.drop(['income', 'native.country', 'hours.per.week', 'capital.loss'], axis=1)
y = df['income']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=42)

categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship']
for feature in categorical:
    le = preprocessing.LabelEncoder()
    X_train[feature] = le.fit_transform(X_train[feature])
    X_test[feature] = le.transform(X_test[feature])

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)

logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred = logreg.predict(X_test)

print('Logistic Regression accuracy score with the first 11 features: {0:0.4f}').
```

Logistic Regression accuracy score with the first 11 features: 0.8186

Select right number of dimensions

```
In [25]: X = df.drop(['income'], axis=1)
y = df['income']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=42)

categorical = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'sex']
for feature in categorical:
    le = preprocessing.LabelEncoder()
    X_train[feature] = le.fit_transform(X_train[feature])
    X_test[feature] = le.transform(X_test[feature])

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)

pca = PCA()
pca.fit(X_train)
cumsum = np.cumsum(pca.explained_variance_ratio_)
dim = np.argmax(cumsum >= 0.90) + 1
print('The number of dimensions required to preserve 90% of variance is', dim)
```


The number of dimensions required to preserve 90% of variance is 12

Plot explained variance ratio with number of dimensions

```
In [28]: plt.figure(figsize=(8,6))
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlim(0,14)
plt.xticks(range(0, 15, 1))
plt.xlabel('Number of components')
plt.ylabel('Cumulative explained variance')
```

Out[28]: Text(0, 0.5, 'Cumulative explained variance')

```
In [29]: plt.show()
```



In []: