

## Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

## Import Dataset

```
In [2]: df = pd.read_csv(r'C:\Users\DELL\Desktop\FSDS\ML\6th - SVM\4th - SVM\SVM\pulsar_
df
```

Out[2]:

	Mean of the integrated profile	Standard deviation of the integrated profile	Excess kurtosis of the integrated profile	Skewness of the integrated profile	Mean of the DM- SNR curve	Standard deviation of the DM-SNR curve	Excess kurtosis of the DM-SNR curve
0	140.562500	55.683782	-0.234571	-0.699648	3.199833	19.110426	7.975532
1	102.507812	58.882430	0.465318	-0.515088	1.677258	14.860146	10.576487
2	103.015625	39.341649	0.323328	1.051164	3.121237	21.744669	7.735822
3	136.750000	57.178449	-0.068415	-0.636238	3.642977	20.959280	6.896499
4	88.726562	40.672225	0.600866	1.123492	1.178930	11.468720	14.269573
...	...	...	...	...	...	...	...
17893	136.429688	59.847421	-0.187846	-0.738123	1.296823	12.166062	15.450260
17894	122.554688	49.485605	0.127978	0.323061	16.409699	44.626893	2.945244
17895	119.335938	59.935939	0.159363	-0.743025	21.430602	58.872000	2.499517
17896	114.507812	53.902400	0.201161	-0.024789	1.946488	13.381731	10.007967
17897	57.062500	85.797340	1.406391	0.089520	188.306020	64.712562	-1.597527

17898 rows × 9 columns



## EDA

```
In [3]: # view dimensions of dataset
df.shape
```

Out[3]: (17898, 9)

```
In [4]: # let's preview the dataset
df.head()
```

Out[4]:

	Mean of the integrated profile	Standard deviation of the integrated profile	Excess kurtosis of the integrated profile	Skewness of the integrated profile	Mean of the DM-SNR curve	Standard deviation of the DM-SNR curve	Excess kurtosis of the DM-SNR curve	Skewness of the DM-SNR curve
0	140.562500	55.683782	-0.234571	-0.699648	3.199833	19.110426	7.975532	74.2
1	102.507812	58.882430	0.465318	-0.515088	1.677258	14.860146	10.576487	127.3
2	103.015625	39.341649	0.323328	1.051164	3.121237	21.744669	7.735822	63.1
3	136.750000	57.178449	-0.068415	-0.636238	3.642977	20.959280	6.896499	53.5
4	88.726562	40.672225	0.600866	1.123492	1.178930	11.468720	14.269573	252.5

In [5]: *# view the column names of the dataframe*

```
col_names = df.columns
col_names
```

```
Out[5]: Index([' Mean of the integrated profile',
              ' Standard deviation of the integrated profile',
              ' Excess kurtosis of the integrated profile',
              ' Skewness of the integrated profile', ' Mean of the DM-SNR curve',
              ' Standard deviation of the DM-SNR curve',
              ' Excess kurtosis of the DM-SNR curve', ' Skewness of the DM-SNR curve',
              'target_class'],
              dtype='object')
```

In [6]: *# remove leading spaces from column names*

```
df.columns = df.columns.str.strip()
```

In [7]: *# view column names again*

```
df.columns
```

```
Out[7]: Index(['Mean of the integrated profile',
              'Standard deviation of the integrated profile',
              'Excess kurtosis of the integrated profile',
              'Skewness of the integrated profile', 'Mean of the DM-SNR curve',
              'Standard deviation of the DM-SNR curve',
              'Excess kurtosis of the DM-SNR curve', 'Skewness of the DM-SNR curve',
              'target_class'],
              dtype='object')
```

In [8]: *# rename column names*

```
df.columns = ['IP Mean', 'IP Sd', 'IP Kurtosis', 'IP Skewness',
              'DM-SNR Mean', 'DM-SNR Sd', 'DM-SNR Kurtosis', 'DM-SNR Skewness',
```

In [9]: *# view the renamed column names*

```
df.columns
```

```
Out[9]: Index(['IP Mean', 'IP Sd', 'IP Kurtosis', 'IP Skewness', 'DM-SNR Mean',
              'DM-SNR Sd', 'DM-SNR Kurtosis', 'DM-SNR Skewness', 'target_class'],
              dtype='object')
```

In [10]: *# check distribution of target\_class column*

```
df['target_class'].value_counts()
```

```
Out[10]: target_class
0      16259
1       1639
Name: count, dtype: int64
```

```
In [11]: # view the percentage distribution of target_class column
df['target_class'].value_counts()/np.float(len(df))
```

```
-----
AttributeError                                Traceback (most recent call last)
Cell In[11], line 2
      1 # view the percentage distribution of target_class column
----> 2 df['target_class'].value_counts()/np.float(len(df))

File ~\anaconda3\Lib\site-packages\numpy\_init_.py:324, in __getattr__(attr)
    319     warnings.warn(
    320         f"In the future `np.{attr}` will be defined as the "
    321         "corresponding NumPy scalar.", FutureWarning, stacklevel=2)
    323 if attr in __former_attrs__:
--> 324     raise AttributeError(__former_attrs__[attr])
    326 if attr == 'testing':
    327     import numpy.testing as testing
```

**AttributeError:** module 'numpy' has no attribute 'float'.  
`np.float` was a deprecated alias for the builtin `float`. To avoid this error in existing code, use `float` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.float64` here. The aliases was originally deprecated in NumPy 1.20; for more details and guidance see the original release note at:  
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
In [12]: df['target_class'].value_counts() / float(len(df))
```

```
Out[12]: target_class
0      0.908426
1      0.091574
Name: count, dtype: float64
```

```
In [13]: # view summary of dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17898 entries, 0 to 17897
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   IP Mean               17898 non-null  float64
 1   IP Sd                 17898 non-null  float64
 2   IP Kurtosis           17898 non-null  float64
 3   IP Skewness           17898 non-null  float64
 4   DM-SNR Mean           17898 non-null  float64
 5   DM-SNR Sd             17898 non-null  float64
 6   DM-SNR Kurtosis       17898 non-null  float64
 7   DM-SNR Skewness       17898 non-null  float64
 8   target_class          17898 non-null  int64
dtypes: float64(8), int64(1)
memory usage: 1.2 MB
```

```
In [14]: # check for missing values in variables
df.isnull().sum()
```

```
Out[14]: IP Mean      0
         IP Sd      0
         IP Kurtosis 0
         IP Skewness 0
         DM-SNR Mean 0
         DM-SNR Sd   0
         DM-SNR Kurtosis 0
         DM-SNR Skewness 0
         target_class 0
         dtype: int64
```

```
In [15]: # view summary statistics in numerical variables
         round(df.describe(),2)
```

```
Out[15]:
```

	IP Mean	IP Sd	IP Kurtosis	IP Skewness	DM-SNR Mean	DM-SNR Sd	DM-SNR Kurtosis	DM-SNR Skewness
<b>count</b>	17898.00	17898.00	17898.00	17898.00	17898.00	17898.00	17898.00	17898.00
<b>mean</b>	111.08	46.55	0.48	1.77	12.61	26.33	8.30	104.86
<b>std</b>	25.65	6.84	1.06	6.17	29.47	19.47	4.51	106.51
<b>min</b>	5.81	24.77	-1.88	-1.79	0.21	7.37	-3.14	-1.98
<b>25%</b>	100.93	42.38	0.03	-0.19	1.92	14.44	5.78	34.96
<b>50%</b>	115.08	46.95	0.22	0.20	2.80	18.46	8.43	83.06
<b>75%</b>	127.09	51.02	0.47	0.93	5.46	28.43	10.70	139.31
<b>max</b>	192.62	98.78	8.07	68.10	223.39	110.64	34.54	1191.00

```
In [16]: # draw boxplots to visualize outliers
```

```
plt.figure(figsize=(24,20))

plt.subplot(4, 2, 1)
fig = df.boxplot(column='IP Mean')
fig.set_title('')
fig.set_ylabel('IP Mean')

plt.subplot(4, 2, 2)
fig = df.boxplot(column='IP Sd')
fig.set_title('')
fig.set_ylabel('IP Sd')

plt.subplot(4, 2, 3)
fig = df.boxplot(column='IP Kurtosis')
fig.set_title('')
fig.set_ylabel('IP Kurtosis')

plt.subplot(4, 2, 4)
fig = df.boxplot(column='IP Skewness')
```

```

fig.set_title('')
fig.set_ylabel('IP Skewness')

plt.subplot(4, 2, 5)
fig = df.boxplot(column='DM-SNR Mean')
fig.set_title('')
fig.set_ylabel('DM-SNR Mean')

plt.subplot(4, 2, 6)
fig = df.boxplot(column='DM-SNR Sd')
fig.set_title('')
fig.set_ylabel('DM-SNR Sd')


plt.subplot(4, 2, 7)
fig = df.boxplot(column='DM-SNR Kurtosis')
fig.set_title('')
fig.set_ylabel('DM-SNR Kurtosis')

plt.subplot(4, 2, 8)
fig = df.boxplot(column='DM-SNR Skewness')
fig.set_title('')
fig.set_ylabel('DM-SNR Skewness')

```

Out[16]: Text(0, 0.5, 'DM-SNR Skewness')

In [17]: plt.show()



```
In [18]: # plot histogram to check distribution

plt.figure(figsize=(24,20))

plt.subplot(4, 2, 1)
fig = df['IP Mean'].hist(bins=20)
fig.set_xlabel('IP Mean')
fig.set_ylabel('Number of pulsar stars')

plt.subplot(4, 2, 2)
fig = df['IP Sd'].hist(bins=20)
fig.set_xlabel('IP Sd')
fig.set_ylabel('Number of pulsar stars')

plt.subplot(4, 2, 3)
fig = df['IP Kurtosis'].hist(bins=20)
fig.set_xlabel('IP Kurtosis')
fig.set_ylabel('Number of pulsar stars')

plt.subplot(4, 2, 4)
fig = df['IP Skewness'].hist(bins=20)
fig.set_xlabel('IP Skewness')
fig.set_ylabel('Number of pulsar stars')

plt.subplot(4, 2, 5)
fig = df['DM-SNR Mean'].hist(bins=20)
fig.set_xlabel('DM-SNR Mean')
fig.set_ylabel('Number of pulsar stars')


plt.subplot(4, 2, 6)
fig = df['DM-SNR Sd'].hist(bins=20)
fig.set_xlabel('DM-SNR Sd')
fig.set_ylabel('Number of pulsar stars')

plt.subplot(4, 2, 7)
fig = df['DM-SNR Kurtosis'].hist(bins=20)
fig.set_xlabel('DM-SNR Kurtosis')
fig.set_ylabel('Number of pulsar stars')

plt.subplot(4, 2, 8)
fig = df['DM-SNR Skewness'].hist(bins=20)
fig.set_xlabel('DM-SNR Skewness')
fig.set_ylabel('Number of pulsar stars')
```

```
Out[18]: Text(0, 0.5, 'Number of pulsar stars')
```

In [19]: `plt.show()`



Declare feature vector and target variable

In [20]: `X = df.drop(['target_class'], axis=1)`  
`y = df['target_class']`

Split data into separate training and test set

In [21]: `# split X and y into training and testing sets`  
`from sklearn.model_selection import train_test_split`  
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)`

In [22]: `# check the shape of X_train and X_test`  
`X_train.shape, X_test.shape`

Out[22]: ((14318, 8), (3580, 8))

Feature Scaling

In [23]: `cols = X_train.columns`

In [24]: `from sklearn.preprocessing import StandardScaler`  
`scaler = StandardScaler()`

```
X_train = scaler.fit_transform(X_train)
```

```
X_test = scaler.transform(X_test)
```

```
In [25]: X_train = pd.DataFrame(X_train, columns=[cols])
```

```
In [26]: X_test = pd.DataFrame(X_test, columns=[cols])
```

```
In [27]: X_train.describe()
```

```
Out[27]:
```

	IP Mean	IP Sd	IP Kurtosis	IP Skewness	DM-SNR Mean	DM
<b>count</b>	1.431800e+04	1.431800e+04	1.431800e+04	1.431800e+04	1.431800e+04	1.431800e+04
<b>mean</b>	1.908113e-16	-6.550610e-16	1.042143e-17	3.870815e-17	-8.734147e-17	-1.600000e+00
<b>std</b>	1.000035e+00	1.000035e+00	1.000035e+00	1.000035e+00	1.000035e+00	1.000035e+00
<b>min</b>	-4.035499e+00	-3.181033e+00	-2.185946e+00	-5.744051e-01	-4.239001e-01	-9.700000e+00
<b>25%</b>	-3.896291e-01	-6.069473e-01	-4.256221e-01	-3.188054e-01	-3.664918e-01	-6.100000e+00
<b>50%</b>	1.587461e-01	5.846646e-02	-2.453172e-01	-2.578142e-01	-3.372294e-01	-4.000000e+00
<b>75%</b>	6.267059e-01	6.501017e-01	-1.001238e-02	-1.419621e-01	-2.463724e-01	1.078000e+00
<b>max</b>	3.151882e+00	7.621116e+00	7.008906e+00	1.054430e+01	7.025568e+00	4.292000e+00

Run SVM with default hyperparameters

```
In [28]: # import SVC classifier
from sklearn.svm import SVC

# import metrics to compute accuracy
from sklearn.metrics import accuracy_score

# instantiate classifier with default hyperparameters
svc=SVC()

# fit classifier to training set
svc.fit(X_train,y_train)

# make predictions on test set
y_pred=svc.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with default hyperparameters: {0:0.4f}'.format(accuracy_score(y_test,y_pred)))
```



Model accuracy score with default hyperparameters: 0.9827

Run SVM with rbf kernel and C=100.0

```
In [29]: # instantiate classifier with rbf kernel and C=100
svc=SVC(C=100.0)

# fit classifier to training set
svc.fit(X_train,y_train)

# make predictions on test set
y_pred=svc.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with rbf kernel and C=100.0 : {0:0.4f}'.format(accu
```

Model accuracy score with rbf kernel and C=100.0 : 0.9832

Run SVM with rbf kernel and C=1000.0

```
In [30]: # instantiate classifier with rbf kernel and C=1000
svc=SVC(C=1000.0)

# fit classifier to training set
svc.fit(X_train,y_train)

# make predictions on test set
y_pred=svc.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with rbf kernel and C=1000.0 : {0:0.4f}'.format(acc
```

Model accuracy score with rbf kernel and C=1000.0 : 0.9816

Run SVM with linear kernel and C=1.0

```
In [31]: # instantiate classifier with linear kernel and C=1.0
linear_svc=SVC(kernel='linear', C=1.0)

# fit classifier to training set
linear_svc.fit(X_train,y_train)

# make predictions on test set
y_pred_test=linear_svc.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with linear kernel and C=1.0 : {0:0.4f}'.format(acc
```

Model accuracy score with linear kernel and C=1.0 : 0.9830

Run SVM with linear kernel and C=100.0

```
In [32]: # instantiate classifier with linear kernel and C=100.0
linear_svc100=SVC(kernel='linear', C=100.0)

# fit classifier to training set
linear_svc100.fit(X_train, y_train)

# make predictions on test set
y_pred=linear_svc100.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with linear kernel and C=100.0 : {0:0.4f}'.format(a
```

Model accuracy score with linear kernel and C=100.0 : 0.9832

Run SVM with linear kernel and C=1000.0

```
In [33]: # instantiate classifier with linear kernel and C=1000.0
linear_svc1000=SVC(kernel='linear', C=1000.0)

# fit classifier to training set
linear_svc1000.fit(X_train, y_train)

# make predictions on test set
y_pred=linear_svc1000.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with linear kernel and C=1000.0 : {0:0.4f}'.format(
```

Model accuracy score with linear kernel and C=1000.0 : 0.9832

Compare the train-set and test-set accuracy

```
In [34]: y_pred_train = linear_svc.predict(X_train)
y_pred_train
```

Out[34]: array([0, 0, 1, ..., 0, 0, 0], dtype=int64)

```
In [35]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_
```

Training-set accuracy score: 0.9783

Check for overfitting and underfitting

```
In [36]: # print the scores on training and test set

print('Training set score: {:.4f}'.format(linear_svc.score(X_train, y_train)))

print('Test set score: {:.4f}'.format(linear_svc.score(X_test, y_test)))
```

Training set score: 0.9783

Test set score: 0.9830

Compare model accuracy with null accuracy

```
In [37]: # check class distribution in test set
y_test.value_counts()
```

```
Out[37]: target_class
0      3306
1       274
Name: count, dtype: int64
```

```
In [38]: # check null accuracy score

null_accuracy = (3306/(3306+274))
print('Null accuracy score: {0:0.4f}'.format(null_accuracy))
```

Null accuracy score: 0.9235

Run SVM with polynomial kernel and C=1.0

```
In [39]: # instantiate classifier with polynomial kernel and C=1.0
poly_svc=SVC(kernel='poly', C=1.0)

# fit classifier to training set
poly_svc.fit(X_train,y_train)

# make predictions on test set
y_pred=poly_svc.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with polynomial kernel and C=1.0 : {0:0.4f}'.format
```

Model accuracy score with polynomial kernel and C=1.0 : 0.9807

Run SVM with polynomial kernel and C=100.0

```
In [40]: # instantiate classifier with polynomial kernel and C=100.0
poly_svc100=SVC(kernel='poly', C=100.0)

# fit classifier to training set
poly_svc100.fit(X_train, y_train)

# make predictions on test set
y_pred=poly_svc100.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with polynomial kernel and C=1.0 : {0:0.4f}'.format
```

Model accuracy score with polynomial kernel and C=1.0 : 0.9824

Run SVM with sigmoid kernel and C=1.0

```
In [41]: # instantiate classifier with sigmoid kernel and C=1.0
sigmoid_svc=SVC(kernel='sigmoid', C=1.0)

# fit classifier to training set
```

```

sigmoid_svc.fit(X_train,y_train)

# make predictions on test set
y_pred=sigmoid_svc.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with sigmoid kernel and C=1.0 : {0:0.4f}'.format(ac

```

Model accuracy score with sigmoid kernel and C=1.0 : 0.8858

Run SVM with sigmoid kernel and C=100.0

```

In [42]: # instantiate classifier with sigmoid kernel and C=100.0
sigmoid_svc100=SVC(kernel='sigmoid', C=100.0)

# fit classifier to training set
sigmoid_svc100.fit(X_train,y_train)

# make predictions on test set
y_pred=sigmoid_svc100.predict(X_test)

# compute and print accuracy score
print('Model accuracy score with sigmoid kernel and C=100.0 : {0:0.4f}'.format(

```

Model accuracy score with sigmoid kernel and C=100.0 : 0.8855

Confusion matrix

```

In [43]: # Print the Confusion Matrix and slice it into four pieces

from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred_test)

print('Confusion matrix\n\n', cm)

print('\nTrue Positives(TP) = ', cm[0,0])

print('\nTrue Negatives(TN) = ', cm[1,1])

print('\nFalse Positives(FP) = ', cm[0,1])

print('\nFalse Negatives(FN) = ', cm[1,0])

```

Confusion matrix

```
[[3289  17]
 [ 44 230]]
```

True Positives(TP) = 3289

True Negatives(TN) = 230

False Positives(FP) = 17

False Negatives(FN) = 44


```
In [44]: # visualize confusion matrix with seaborn heatmap

cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                        index=['Predict Positive:1', 'Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

Out[44]: <Axes: >

```
In [45]: plt.show()
```



Classification Report

```
In [46]: from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	3306
1	0.93	0.84	0.88	274
accuracy			0.98	3580
macro avg	0.96	0.92	0.94	3580
weighted avg	0.98	0.98	0.98	3580

Classification accuracy

```
In [47]: TP = cm[0,0]
        TN = cm[1,1]
        FP = cm[0,1]
        FN = cm[1,0]
```

```
In [48]: # print classification accuracy

classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)

print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))
```

Classification accuracy : 0.9830

Classification error

```
In [49]: # print classification error

classification_error = (FP + FN) / float(TP + TN + FP + FN)

print('Classification error : {0:0.4f}'.format(classification_error))
```

Classification error : 0.0170

Precision

```
In [50]: # print precision score

precision = TP / float(TP + FP)

print('Precision : {0:0.4f}'.format(precision))
```

Precision : 0.9949

Recall

```
In [51]: recall = TP / float(TP + FN)

print('Recall or Sensitivity : {0:0.4f}'.format(recall))
```

Recall or Sensitivity : 0.9868

True Positive Rate

```
In [52]: true_positive_rate = TP / float(TP + FN)

print('True Positive Rate : {0:0.4f}'.format(true_positive_rate))
```

True Positive Rate : 0.9868

False Positive Rate

```
In [53]: false_positive_rate = FP / float(FP + TN)

print('False Positive Rate : {0:0.4f}'.format(false_positive_rate))
```

False Positive Rate : 0.0688

Specificity

```
In [54]: specificity = TN / (TN + FP)
print('Specificity : {0:0.4f}'.format(specificity))
```

Specificity : 0.9312

ROC - AUC

```
In [55]: # plot ROC Curve

from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_test, y_pred_test)

plt.figure(figsize=(6,4))

plt.plot(fpr, tpr, linewidth=2)

plt.plot([0,1], [0,1], 'k--' )


plt.rcParams['font.size'] = 12

plt.title('ROC curve for Predicting a Pulsar Star classifier')

plt.xlabel('False Positive Rate (1 - Specificity)')

plt.ylabel('True Positive Rate (Sensitivity)')

plt.show()
```



```
In [56]: # compute ROC AUC

from sklearn.metrics import roc_auc_score

ROC_AUC = roc_auc_score(y_test, y_pred_test)

print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

ROC AUC : 0.9171

```
In [57]: # calculate cross-validated ROC AUC

from sklearn.model_selection import cross_val_score

Cross_validated_ROC_AUC = cross_val_score(linear_svc, X_train, y_train, cv=10, s

print('Cross validated ROC AUC : {:.4f}'.format(Cross_validated_ROC_AUC))
```

Cross validated ROC AUC : 0.9756

Stratified k-Fold Cross Validation with shuffle split with linear kernel

```
In [58]: from sklearn.model_selection import KFold

kfold=KFold(n_splits=5, shuffle=True, random_state=0)

linear_svc=SVC(kernel='linear')

linear_scores = cross_val_score(linear_svc, X, y, cv=kfold)
```

```
In [59]: # print cross-validation scores with linear kernel
```



```
print('Stratified cross-validation scores with linear kernel:\n\n{}'.format(line
```

Stratified cross-validation scores with linear kernel:

```
[0.98296089 0.97458101 0.97988827 0.97876502 0.97848561]
```

In [60]: *# print average cross-validation score with linear kernel*

```
print('Average stratified cross-validation score with linear kernel:{:.4f}'.form
```

Average stratified cross-validation score with linear kernel:0.9789

Stratified k-Fold Cross Validation with shuffle split with rbf kernel

In [61]: `rbf_svc=SVC(kernel='rbf')`

```
rbf_scores = cross_val_score(rbf_svc, X, y, cv=kfold)
```

In [62]: *# print cross-validation scores with rbf kernel*

```
print('Stratified Cross-validation scores with rbf kernel:\n\n{}'.format(rbf_sco
```

Stratified Cross-validation scores with rbf kernel:

```
[0.97849162 0.97011173 0.97318436 0.9709416 0.96982397]
```

In [63]: *# print average cross-validation score with rbf kernel*

```
print('Average stratified cross-validation score with rbf kernel:{:.4f}'.format(
```

Average stratified cross-validation score with rbf kernel:0.9725

Hyperparameter Optimization using GridSearch CV

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

# 1. Prepare your data
# Assume df is already defined and cleaned
X = df.drop('target_class', axis=1) # Features
y = df['target_class']             # Target

# 2. Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# 3. Instantiate base SVC model
svc = SVC()

# 4. Declare parameter grid
parameters = [
    {'C': [1, 10, 100, 1000], 'kernel': ['linear']},
    {'C': [1, 10, 100, 1000], 'kernel': ['rbf'],
     'gamma': [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]},
    {'C': [1, 10, 100, 1000], 'kernel': ['poly'],
     'degree': [2, 3, 4], 'gamma': [0.01, 0.02, 0.03, 0.04, 0.05]}
]
```

```

# 5. GridSearchCV
grid_search = GridSearchCV(
    estimator=svc,
    param_grid=parameters,
    scoring='accuracy',
    cv=5,
    verbose=0
)

# 6. Fit to training data
grid_search.fit(X_train, y_train)

# 7. Best parameters and score
print("Best Parameters:", grid_search.best_params_)
print("Best Cross-Validation Score:", grid_search.best_score_)

# 8. Evaluate on test set
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
print("Test Accuracy:", accuracy_score(y_test, y_pred))

```

```

In [ ]: # examine the best model

# best score achieved during the GridSearchCV
print('GridSearch CV best score : {:.4f}\n\n'.format(grid_search.best_score_))

# print parameters that give the best results
print('Parameters that give the best results :','\n\n', (grid_search.best_params_))

# print estimator that was chosen by the GridSearch
print('\n\nEstimator that was chosen by the search :','\n\n', (grid_search.best_estimator_))

```

```

In [ ]: # calculate GridSearch CV score on test set

print('GridSearch CV score on test set: {0:0.4f}'.format(grid_search.score(X_test, y_test)))

```

```

In [ ]:

```